



AMERICAN INTERNATIONAL UNIVERSITY – BANGLADESH

INTRODUCTION TO DATA SCIENCE

FALL- 2024-2025

Supervise By

Tohedul Islam

Assignment: Mid Term Project Report

Group No: 02

Section: C

Submitted By:

NAME	ID
FARJANA YESMIN OPI	22-47018-1
MD. ABU TOWSIF	22-47019-1
MOST. SAYMA KHATUN	22-47035-1

TABLE OF CONTENTS

Topics	Page no.
1. Description of the dataset	03
2. Loading the dataset	03-04
3. Dataset dimensions and structures	04-05
5. Extracting unique values from dataset columns	05
6. Removing duplicate rows from the dataset	06
7. Handling invalid values	06-08
8. Dealing with Missing Values	09-13
9. Converting Categorical Columns to Numeric Factors	13-14
10. Identifying Outliers	14-15
11. Removing Outliers	15-16
12. Normalizing the Dataset	17-18
13. Displaying Summary of the	18-19
14. Measure of Central Tendency	19-20
15. Measure of Spread	20-22
16. Handling Imbalance dataset	23-24

Description of the dataset

The Loan Approval Classification dataset is a synthetic dataset created for binary classification tasks focused on loan approval decisions. It contains information about individuals applying for loans, described by 14 attributes. These attributes provide a comprehensive profile of each applicant and their loan application details. The dataset includes the applicant's age (person_age), gender (person_gender), highest education level (person_education), annual income (person_income), and years of employment experience (person_emp_exp). Additionally, it captures the applicant's home ownership status (person_home_ownership) and their creditworthiness through the credit history length (cb_person_cred_hist_length) and credit score (credit_score). The dataset also includes information specific to the loan application, such as the loan amount requested (loan_amnt), the purpose of the loan (loan_intent), the loan interest rate (loan_int_rate), and the loan amount as a percentage of the applicant's annual income (loan_percent_income). An indicator of previous loan defaults (previous_loan_defaults_on_file) is also provided. The target variable, loan_status, indicates the final approval status of the loan, where 1 represents approval and 0 represents rejection.

This dataset contains both numerical and categorical data, offering a rich set of features for exploring and modeling factors that influence loan approval decisions.

Loading the dataset

Code

```
install.packages("dplyr")  
library(dplyr)  
install.packages("readxl");  
library(readxl);  
dataSet_1 <- read_excel("E:/FALL 24-25/INTRODUCTION TO DATA SCIENCE/MID/Mid  
Term Project/Materials/Midterm_Dataset_Section(C).xlsx");  
print(dataSet_1, n = nrow(dataSet_1));
```

Output

```
R - R 4.2.1 - E:/FALL 24-25/INTRODUCTION TO DATA SCIENCE/MID/Mid Term Project/ > library(dplyr) > library(readxl) > dataSet_1 <- read_excel("E:/FALL 24-25/INTRODUCTION TO DATA SCIENCE/MID/Mid Term Project/Metaterials/Midterm_Dataset_Section(C).xlsx"); > print(dataSet_1, n = nrow(dataSet_1)); # A tibble: 201 x 14
```

	person_age	person_gender	person_education	person_income	person_emp_exp	person_home_ownership	loan_amnt	loan_intent	loan_int_rate
1	21	female	Master	71948	0	RENT	35000	PERSONAL	16.0
2	21	female	High School	12282	0	OWN	1000	EDUCATION	11.1
3	25	female	High School	12438	3	MORTGAGE	5500	MEDICAL	12.9
4	23	female	Bachelor	79753	0	RENT	35000	MEDICAL	15.2
5	24	male	Master	66135	1	RENTT	35000	MEDICAL	14.3
6	NA	female	High School	12951	0	OWN	2500	VENTURE	7.14
7	22	female	Bachelor	NA	1	RENT	35000	EDUCATION	12.4
8	24	NA	High School	95550	5	RENT	35000	MEDICAL	11.1
9	22	female	NA	100684	3	RENT	35000	PERSONAL	8.9
10	21	female	High School	12739	0	OWN	1600	VENTURE	14.7
11	22	female	High School	102985	0	RENT	35000	VENTURE	10.4
12	21	female	Associate	13113	0	OWN	4500	HOMEIMPROVEMENT	8.63
13	23	male	Bachelor	114860	3	RENT	35000	VENTURE	7.9
14	NA	male	Master	130713	0	RENT	35000	EDUCATION	18.4
15	23	female	Associate	3138998	0	RENT	35000	EDUCATION	7.9
16	23	female	NA	NA	5	MORTGAGE	30000	DEBTCONSOLIDATION	10.6
17	23	NA	Bachelor	144943	0	RENT	35000	EDUCATION	7.9
18	23	female	High School	111369	0	RENT	35000	MEDICAL	20
19	23	male	Bachelor	136628	0	RENT	35000	DEBTCONSOLIDATION	18.2
20	24	female	Master	14283	1	MORTGAGE	1750	EDUCATION	11.0
21	25	male	Bachelor	195718	0	RENT	35000	VENTURE	7.49
22	25	male	High School	165792	4	RENT	34800	PERSONAL	16.8
23	22	female	Master	79255	0	RENT	34000	EDUCATION	17.6
24	24	female	Bachelor	13866	0	OWN	1500	PERSONAL	7.29

Description

The dplyr package is installed and loaded for data manipulation, while the readxl package is installed and loaded for reading Excel files. The dataset Midterm_Dataset_Section(C).xlsx is read from the specified file path into dataSet_1. Finally, the entire dataset is printed using print with all rows displayed. In the screenshot above, we can see the first 24 instances of the dataset. Though the output displayed all the instances of the dataset.

Dataset Dimensions and Structure

Code

```
no_of_col <- ncol(dataSet_1)

no_of_row <- nrow(dataSet_1)

cat("No of row in the dataset: ", no_of_row)

cat("No of column in the dataset: ", no_of_col)

str(dataSet_1)
```

Output

```
> no_of_col <- ncol(dataSet_1)
> no_of_row <- nrow(dataSet_1)
> cat("No of row in the dataset: ", no_of_row)
No of row in the dataset: 201
> cat("No of column in the dataset: ", no_of_col)
No of column in the dataset: 14
> str(dataSet_1)
tibble [201 × 14] (S3: tbl_df/tbl/data.frame)
 $ person_age      : num [1:201] 21 21 25 23 24 NA 22 24 22 21 ...
 $ person_gender   : chr [1:201] "female" "female" "female" "female" ...
 $ person_education : chr [1:201] "Master" "High School" "High School" "Bachelor" ...
 $ person_income   : num [1:201] 71948 12282 12438 79753 66135 ...
 $ person_emp_exp   : num [1:201] 0 0 3 0 1 0 1 5 3 0 ...
 $ person_home_ownership : chr [1:201] "RENT" "OWN" "MORTGAGE" "RENT" ...
 $ loan_amnt       : num [1:201] 35000 1000 5500 35000 35000 2500 35000 35000 1600 ...
 $ loan_intent      : chr [1:201] "PERSONAL" "EDUCATION" "MEDICAL" "MEDICAL" ...
 $ loan_int_rate    : num [1:201] 16 11.1 12.9 15.2 14.3 ...
 $ loan_percent_income : num [1:201] 0.49 NA 0 0.44 0.53 0.19 0.37 0.37 0.35 0.13 ...
 $ cb_person_cred_hist_length : num [1:201] 3 2 3 2 4 2 3 4 2 3 ...
 $ credit_score     : num [1:201] 561 504 635 675 586 532 701 585 544 640 ...
 $ previous_loan_defaults_on_file : chr [1:201] "No" "Yes" "No" "No" ...
 $ loan_status      : num [1:201] 1 0 1 1 1 1 1 1 NA 1 ...
> |
```

Description

This code calculates and prints the number of rows and columns in the dataset `dataSet_1` using the `nrow` and `ncol` functions. Additionally, the `str` function provides a detailed overview of the dataset's structure, including the column names, data types, and sample values for each column. The output includes both the dataset's dimensions and a concise summary of its attributes.

Extracting Unique Values from Dataset Columns

Code

```
unique(dataSet_1$person_age)

unique(dataSet_1$person_gender)
```

Output

```
> unique(dataSet_1$person_age)
[1] 21 25 23 24 NA 22 230 26 350 144
> unique(dataSet_1$person_gender)
[1] "female" "male" NA
> |
```

Description

We are extracting unique values from specific columns in the dataset `dataSet_1`. By applying the `unique` function to the `person_age` column, we retrieve all distinct age values, while applying it to the `person_gender` column provides a list of unique gender categories. We can see NULL values in `person_gender` columns, we will deal with them later.

Removing Duplicate Rows from the Dataset

Code

```
remo_dupli_dataset <- distinct(dataSet_1);

remo_dupli_dataset

cat ("No of row and column after removing duplicate instances: ",
nrow(remo_dupli_dataset), ncol(remo_dupli_dataset))
```

Output

```
> remo_dupli_dataset <- distinct(dataSet_1);
> remo_dupli_dataset
# A tibble: 200 x 14
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate
  <dbl> <chr> <chr> <dbl> <dbl> <chr> <dbl> <chr> <dbl>
1 21 female Master 71948 0 RENT 35000 PERSONAL 16.0
2 21 female High School 12282 0 OWN 1000 EDUCATION 11.1
3 25 female High School 12438 3 MORTGAGE 5500 MEDICAL 12.9
4 23 female Bachelor 79753 0 RENT 35000 MEDICAL 15.2
5 24 male Master 66135 1 RENTT 35000 MEDICAL 14.3
6 NA female High School 12951 0 OWN 2500 VENTURE 7.14
7 22 female Bachelor NA 1 RENT 35000 EDUCATION 12.4
8 24 NA High School 95550 5 RENT 35000 MEDICAL 11.1
9 22 female NA 100684 3 RENT 35000 PERSONAL 8.9
10 21 female High School 12739 0 OWN 1600 VENTURE 14.7
# i 190 more rows
# i 5 more variables: loan_percent_income <dbl>, cb_person_cred_hist_length <dbl>, credit_score <dbl>, previous_loan_defaults_on_file <chr>,
# loan_status <dbl>
# i Use `print(n = ...)` to see more rows
> cat ("No of row and column after removing duplicate instances: ", nrow(remo_dupli_dataset), ncol(remo_dupli_dataset))
No of row and column after removing duplicate instances: 200 14
>|
```

Description

We are removing duplicate rows from the dataset dataSet_1 to ensure each entry is unique. The resulting dataset, stored in remo_dupli_dataset, is displayed, followed by the total number of rows and columns remaining after duplicate removal. One instance has been removed, the the dimension of the new dataset is 200 rows and 14 columns

Handling Invalid Values

Code

```
fresh_dataset <- remo_dupli_dataset;

unique(fresh_dataset$person_gender)

unique(fresh_dataset$person_education)

unique(fresh_dataset$person_home_ownership)

unique(fresh_dataset$loan_intent)
```

```

unique(fresh_dataset$previous_loan_defaults_on_file);

fresh_dataset$person_age[is.na(as.numeric(as.character(fresh_dataset$person_age)))]
fresh_dataset$person_income[is.na(as.numeric(as.character(fresh_dataset$person_income))
)]
fresh_dataset$person_emp_exp[is.na(as.numeric(as.character(fresh_dataset$person_emp_exp
)))]
fresh_dataset$loan_amnt[is.na(as.numeric(as.character(fresh_dataset$loan_amnt)))]
fresh_dataset$loan_int_rate[is.na(as.numeric(as.character(fresh_dataset$loan_int_rate))
)]
fresh_dataset$loan_percent_income[is.na(as.numeric(as.character(fresh_dataset$loan_perc
ent_income)))]
fresh_dataset$cb_person_cred_hist_length[is.na(as.numeric(as.character(fresh_dataset$cb
_person_cred_hist_length)))]
fresh_dataset$credit_score[is.na(as.numeric(as.character(fresh_dataset$credit_score)))]
fresh_dataset$loan_status[is.na(as.numeric(as.character(fresh_dataset$loan_status)))]

deal_invalid_dataset <- fresh_dataset;
deal_invalid_dataset$person_home_ownership <- ifelse(
  substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 2) == "OT", "OTHER",
  ifelse(
    substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 1) == "O", "OWN",
    ifelse(
      substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 1) == "R", "RENT",
      ifelse(
        substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 1) == "M", "MORTGAGE",
        "NA"
      )
    )
  )
)
)

```

Output

```
> fresh_dataset <- remo_dupli_dataset;
> unique(fresh_dataset$person_gender)
[1] "female" "male" NA
> unique(fresh_dataset$person_education)
[1] "Master" "High School" "Bachelor" NA "Associate" "Doctorate"
> unique(fresh_dataset$person_home_ownership)
[1] "RENT" "OWN" "MORTGAGE" "RENTT" "OOWN" "OTHER"
> unique(fresh_dataset$loan_intent)
[1] "PERSONAL" "EDUCATION" "MEDICAL" "VENTURE" "HOMEIMPROVEMENT" "DEBTCONSOLIDATION"
> unique(fresh_dataset$previous_loan_defaults_on_file);
[1] "No" "Yes"
> fresh_dataset$person_age[is.na(as.numeric(as.character(fresh_dataset$person_age)))]
[1] NA NA NA NA
> fresh_dataset$person_income[is.na(as.numeric(as.character(fresh_dataset$person_income)))]
[1] NA NA NA NA
> fresh_dataset$person_emp_exp[is.na(as.numeric(as.character(fresh_dataset$person_emp_exp)))]
numeric(0)
> fresh_dataset$loan_amnt[is.na(as.numeric(as.character(fresh_dataset$loan_amnt)))]
numeric(0)
> fresh_dataset$loan_int_rate[is.na(as.numeric(as.character(fresh_dataset$loan_int_rate)))]
numeric(0)
> fresh_dataset$loan_percent_income[is.na(as.numeric(as.character(fresh_dataset$loan_percent_income)))]
[1] NA
> fresh_dataset$cb_person_cred_hist_length[is.na(as.numeric(as.character(fresh_dataset$cb_person_cred_hist_length)))]
numeric(0)
> fresh_dataset$credit_score[is.na(as.numeric(as.character(fresh_dataset$credit_score)))]
numeric(0)
> fresh_dataset$loan_status[is.na(as.numeric(as.character(fresh_dataset$loan_status)))]
[1] NA NA NA
> deal_invalid_dataset <- fresh_dataset;
> deal_invalid_dataset$person_home_ownership <- ifelse(
+   substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 2) == "OT", "OTHER",
+   ifelse(
+     substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 1) == "O", "OWN",
+     ifelse(
+       substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 1) == "R", "RENT",
+       ifelse(
+         substr(toupper(deal_invalid_dataset$person_home_ownership), 1, 1) == "M", "MORTGAGE",
+         "NA"
+       )
+     )
+   )
+ )
> unique(deal_invalid_dataset $person_home_ownership)
[1] "RENT" "OWN" "MORTGAGE" "OTHER"
> |
```

Description

We first checked the unique values of categorical columns in the dataset, such as `person_gender`, `person_education`, `person_home_ownership`, `loan_intent`, and `previous_loan_defaults_on_file`. Next, we verified whether numerical columns contained only valid numeric values, and identified any missing values (NAs) which would be addressed later. We then focused on cleaning the `person_home_ownership` column, where invalid values were present. We assumed that if the value starts with "OT," it should be categorized as "OTHER." Similarly, values starting with "O" were classified as "OWN," those starting with "R" as "RENT," those starting with "M" as "MORTGAGE," and any other value was assigned "NA." Finally, we confirmed the changes by examining the unique values in each column using the `unique()` function.

Dealing with Missing Values

Discard instances

Code

```
fresh_dataset <- deal_invalid_dataset;
deal_miss_value_dataset <- fresh_dataset;
colSums(is.na(deal_miss_value_dataset));

which(is.na(deal_miss_value_dataset$ person_age))

deal_miss_value_dataset <- na.omit(deal_miss_value_dataset);
colSums(is.na(deal_miss_value_dataset));
```

Output

```
> fresh_dataset <- deal_invalid_dataset;
> deal_miss_value_dataset <- fresh_dataset;
> colSums(is.na(deal_miss_value_dataset));
      person_age      person_gender      person_education      person_income
           4              4              2              4
      person_emp_exp      person_home_ownership      loan_amnt      loan_intent
           0              0              0              0
      loan_int_rate      loan_percent_income      cb_person_cred_hist_length      credit_score
           0              1              0              0
previous_loan_defaults_on_file      loan_status
           0              3
> which(is.na(deal_miss_value_dataset$ person_age))
[1]  6 14 28 34
> deal_miss_value_dataset <- na.omit(deal_miss_value_dataset);
> colSums(is.na(deal_miss_value_dataset));
      person_age      person_gender      person_education      person_income
           0              0              0              0
      person_emp_exp      person_home_ownership      loan_amnt      loan_intent
           0              0              0              0
      loan_int_rate      loan_percent_income      cb_person_cred_hist_length      credit_score
           0              0              0              0
previous_loan_defaults_on_file      loan_status
           0              0
> |
```

Description

We began by examining the missing values in the dataset by using `colSums(is.na())` to get a summary of missing values across all columns. The number of missing values in each column is displayed. Then, we identified specific rows where the `person_age` column contained missing values with the `which(is.na())` function. To address these missing values, we removed any rows containing NA values from the dataset using the `na.omit()` function. Finally, we reassessed the dataset to confirm that all missing values were successfully removed by applying `colSums(is.na())` once more. We can see no column contains any missing values

Top-Down and Bottom-Up Approach

Code

```
top_down_dataset <- fresh_dataset %>% fill(person_age, person_gender,
person_education, person_income, loan_percent_income, loan_status, .direction = 'down')

colSums(is.na(top_down_dataset));

bottom_up_dataset <- fresh_dataset %>% fill(person_age, person_gender, person_education,
person_income, loan_percent_income, loan_status, .direction = 'up')

colSums(is.na(bottom_up_dataset));
```

Output

```
> top_down_dataset <- fresh_dataset %>% fill(person_age, person_gender, person_education, person_income, loan_percent_income, loan_status, .direction =
'down')
> colSums(is.na(top_down_dataset));
      person_age      person_gender      person_education      person_income
      person_emp_exp      person_home_ownership      loan_amnt      loan_intent
      loan_int_rate      loan_percent_income      cb_person_cred_hist_length      credit_score
previous_loan_defaults_on_file      loan_status
      0
> bottom_up_dataset <- fresh_dataset %>% fill(person_age, person_gender, person_education, person_income, loan_percent_income, loan_status, .direction =
'up')
> colSums(is.na(bottom_up_dataset));
      person_age      person_gender      person_education      person_income
      person_emp_exp      person_home_ownership      loan_amnt      loan_intent
      loan_int_rate      loan_percent_income      cb_person_cred_hist_length      credit_score
previous_loan_defaults_on_file      loan_status
      0
>
> |
```

Description

We applied two approaches to fill missing values in the dataset. In the Top-Down approach, we used the fill() function with the .direction = 'down' parameter to fill missing values by propagating the previous value downward across selected columns. We then checked for any remaining missing values using colSums(is.na()).

In the Bottom-Up approach, we again used the fill() function but with the .direction = 'up' parameter, which propagates missing values upward. We confirmed the absence of any remaining missing values by examining the result with colSums(is.na()).

Replace by Most Frequent/Average Value

For categorical columns (Mode)

Code

```
deal_miss_value_mode <- fresh_dataset;

mode_person_gender <- names(sort(table(deal_miss_value_mode$person_gender), decreasing
= TRUE))[1]

deal_miss_value_mode$person_gender[is.na(deal_miss_value_mode$person_gender)] <-
mode_person_gender

mode_person_education <- names(sort(table(deal_miss_value_mode$person_education),
decreasing = TRUE))[1]

deal_miss_value_mode$person_education[is.na(deal_miss_value_mode$person_education)] <-
mode_person_education

mode_person_home_ownership <-
names(sort(table(deal_miss_value_mode$person_home_ownership), decreasing = TRUE))[1]

deal_miss_value_mode$person_home_ownership[is.na(deal_miss_value_mode$person_home_ownership)] <- mode_person_home_ownership

mode_loan_intent <- names(sort(table(deal_miss_value_mode$loan_intent), decreasing =
TRUE))[1]

deal_miss_value_mode$loan_intent[is.na(deal_miss_value_mode$loan_intent)] <-
mode_loan_intent

mode_previous_loan_defaults_on_file <-
names(sort(table(deal_miss_value_mode$previous_loan_defaults_on_file), decreasing =
TRUE))[1]

deal_miss_value_mode$previous_loan_defaults_on_file[is.na(deal_miss_value_mode$previous
_loan_defaults_on_file)] <- mode_previous_loan_defaults_on_file

colSums(is.na(deal_miss_value_mode))
```

Output

```
> deal_miss_value_mode <- fresh_dataset;
> mode_person_gender <- names(sort(table(deal_miss_value_mode$person_gender), decreasing = TRUE))[1]
> deal_miss_value_mode$person_gender[is.na(deal_miss_value_mode$person_gender)] <- mode_person_gender
> mode_person_education <- names(sort(table(deal_miss_value_mode$person_education), decreasing = TRUE))[1]
> deal_miss_value_mode$person_education[is.na(deal_miss_value_mode$person_education)] <- mode_person_education
> mode_person_home_ownership <- names(sort(table(deal_miss_value_mode$person_home_ownership), decreasing = TRUE))[1]
> deal_miss_value_mode$person_home_ownership[is.na(deal_miss_value_mode$person_home_ownership)] <- mode_person_home_ownership
> mode_loan_intent <- names(sort(table(deal_miss_value_mode$loan_intent), decreasing = TRUE))[1]
> deal_miss_value_mode$loan_intent[is.na(deal_miss_value_mode$loan_intent)] <- mode_loan_intent
> mode_previous_loan_defaults_on_file <- names(sort(table(deal_miss_value_mode$previous_loan_defaults_on_file), decreasing = TRUE))[1]
> deal_miss_value_mode$previous_loan_defaults_on_file[is.na(deal_miss_value_mode$previous_loan_defaults_on_file)] <- mode_previous_loan_defaults_on_file
> colSums(is.na(deal_miss_value_mode))
  person_age      person_gender      person_education      person_income
         4                0                0                4
  person_emp_exp  person_home_ownership      loan_amnt      loan_intent
         0                0                0                0
  loan_int_rate  loan_percent_income  cb_person_cred_hist_length  credit_score
         0                1                0                0
previous_loan_defaults_on_file      loan_status
         0                3
> |
```

Description

We handled missing values in categorical columns by replacing them with the most frequent value (mode). For each categorical column—`person_gender`, `person_education`, `person_home_ownership`, `loan_intent`, and `previous_loan_defaults_on_file`—we first identified the mode using the `sort()` and `table()` functions. Then, we replaced any missing values with these most common values. Finally, we checked if any missing values remained in the dataset by summarizing with `colSums(is.na())`.

For numerical columns (mean)

Code

```
deal_miss_value_mean <- deal_miss_value_mode;

for(col_name in c("person_age", "person_income", "loan_percent_income", "loan_status")) {

  if(is.numeric(deal_miss_value_mean[[col_name]])) {

    column_mean <- mean(deal_miss_value_mean[[col_name]], na.rm = TRUE)

    deal_miss_value_mean[[col_name]][is.na(deal_miss_value_mean[[col_name]])] <- column_mean

    deal_miss_value_mean[[col_name]] <- round(deal_miss_value_mean[[col_name]], digits = 0)

  }

}

colSums(is.na(deal_miss_value_mean))
```

Output

```
> deal_miss_value_mean <- deal_miss_value_mode;
> for(col_name in c("person_age", "person_income", "loan_percent_income", "loan_status")) {
+   if(is.numeric(deal_miss_value_mean[[col_name]])) {
+     column_mean <- mean(deal_miss_value_mean[[col_name]], na.rm = TRUE)
+     deal_miss_value_mean[[col_name]][is.na(deal_miss_value_mean[[col_name]])] <- column_mean
+     deal_miss_value_mean[[col_name]] <- round(deal_miss_value_mean[[col_name]], digits = 0)
+   }
+ }
> colSums(is.na(deal_miss_value_mean))
      person_age      person_gender      person_education      person_income
              0              0              0              0
  person_emp_exp  person_home_ownership      loan_amnt      loan_intent
              0              0              0              0
    loan_int_rate  loan_percent_income  cb_person_cred_hist_length      credit_score
              0              0              0              0
previous_loan_defaults_on_file      loan_status
              0              0
> |
```

Description

We replaced missing values in numerical columns—`person_age`, `person_income`, `loan_percent_income`, and `loan_status`—by using the mean value of each respective column. For

each column, we calculated the mean while excluding missing values, rounded the result to the nearest integer, and substituted any missing entries with this mean. Finally, we checked if any missing values remained in the dataset using `colSums(is.na())`.

Converting Categorical Columns to Numeric Factors

Code

```
fresh_dataset <- deal_miss_value_dataset;

dataSet_num <- fresh_dataset;

dataSet_num$person_gender <- factor(dataSet_num$person_gender, levels = c("male",
"female"), labels = c(1,2));

dataSet_num$person_education <- factor(dataSet_num$person_education, levels = c("High
School", "Bachelor", "Master", "Associate", "Doctorate"), labels = c(1,2,3,4,5));

dataSet_num$loan_intent <- factor(dataSet_num$loan_intent, levels =
c("PERSONAL","EDUCATION","MEDICAL","VENTURE","HOMEIMPROVEMENT", "DEBTCONSOLIDATION"),
labels = c(1,2,3,4,5, 6));

dataSet_num$person_home_ownership <- factor(dataSet_num$person_home_ownership, levels
= c("RENT","OWN","MORTGAGE","OTHER"), labels = c(1,2,3,4));

dataSet_num$previous_loan_defaults_on_file <-
factor(dataSet_num$previous_loan_defaults_on_file, levels = c("Yes", "No"), labels =
c(1,2));

dataSet_num
```

Output

```
> fresh_dataset <- deal_miss_value_dataset;
> dataSet_num <- fresh_dataset;
> dataSet_num$person_gender <- factor(dataSet_num$person_gender, levels = c("male", "female"), labels = c(1,2));
> dataSet_num$person_education <- factor(dataSet_num$person_education, levels = c("High School", "Bachelor", "Master", "Associate", "Doctorate"), labels = c(1,2,3,4,5));
> dataSet_num$loan_intent <- factor(dataSet_num$loan_intent, levels = c("PERSONAL","EDUCATION","MEDICAL","VENTURE","HOMEIMPROVEMENT", "DEBTCONSOLIDATION"), labels = c(1,2,3,4,5, 6));
> dataSet_num$person_home_ownership <- factor(dataSet_num$person_home_ownership, levels = c("RENT","OWN","MORTGAGE","OTHER"), labels = c(1,2,3,4));
> dataSet_num$previous_loan_defaults_on_file <- factor(dataSet_num$previous_loan_defaults_on_file, levels = c("Yes", "No"), labels = c(1,2));
> dataSet_num
# A tibble: 184 x 14
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate
  <dbl> <fct> <fct> <dbl> <dbl> <fct> <dbl> <fct> <dbl>
1 21 2 3 71948 0 1 35000 1 16.0
2 25 2 1 12438 3 3 5500 3 12.9
3 23 2 2 79753 0 1 35000 3 15.2
4 24 1 3 66135 1 1 35000 3 14.3
5 21 2 1 12739 0 2 1600 4 14.7
6 22 2 1 102985 0 1 35000 4 10.4
7 21 2 4 13113 0 2 4500 5 8.63
8 23 1 2 114860 3 1 35000 4 7.9
9 23 1 2 136628 0 1 35000 6 18.2
10 24 2 3 14283 1 3 1750 2 11.0
# i 174 more rows
# i 5 more variables: loan_percent_income <dbl>, cb_person_cred_hist_length <dbl>, credit_score <dbl>, previous_loan_defaults_on_file <fct>,
# loan_status <dbl>
# i Use 'print(n = ...)' to see more rows
> |
```

Description

We deal the missing values with 3 methods. But we will use the dataset we got after discarding missing values for further analysis. We converted all categorical columns into numeric factors to prepare the dataset for analysis. The conversion included the following mappings: person_gender was encoded as 1 (male) and 2 (female), person_education was mapped across five levels (High School, Associate, Bachelor, Master, Doctorate), loan_intent was mapped to six loan purposes (Personal, Education, Medical, Venture, Home Improvement, Debt Consolidation), person_home_ownership was categorized into four types (Rent, Own, Mortgage, Other), and previous_loan_defaults_on_file was encoded as 1 (Yes) and 2 (No).

Identifying Outliers

Code

```
detect_outlier <- function(dataframe, columns) {
  for (col in columns) {
    if (is.numeric(dataframe[[col]])) {
      Quantile1 <- quantile(dataframe[[col]], probs = 0.25)
      Quantile3 <- quantile(dataframe[[col]], probs = 0.75)
      IQR <- Quantile3 - Quantile1
      outlier_flags <- dataframe[[col]] > Quantile3 + (IQR * 1.5) |
        dataframe[[col]] < Quantile1 - (IQR * 1.5)
      outliers <- dataframe[[col]][outlier_flags]
      if (length(outliers) > 0) {
        cat("Outliers detected in column", col, ":\n")
        print(outliers)
      } else {
        cat("No outliers detected in column", col, "\n")
      }
    } else {
      cat("Column", col, "is not numeric, skipped\n")
    }
  }
}

detect_outlier(fresh_dataset, names(fresh_dataset))
```

Output

```
> detect_outlier(fresh_dataset, names(fresh_dataset))
Outliers detected in column person_age :
[1] 230 350 144 144
Column person_gender is not numeric, skipped
Column person_education is not numeric, skipped
No outliers detected in column person_income
Outliers detected in column person_emp_exp :
[1] 125 8 121
Column person_home_ownership is not numeric, skipped
No outliers detected in column loan_amnt
Column loan_intent is not numeric, skipped
Outliers detected in column loan_int_rate :
[1] 5.42 19.91
No outliers detected in column loan_percent_income
No outliers detected in column cb_person_cred_hist_length
Outliers detected in column credit_score :
[1] 789 484 807
Column previous_loan_defaults_on_file is not numeric, skipped
No outliers detected in column loan_status
> |
```

Description

We applied a user defined `detect_outlier` function to identify outlier values in each numeric column of the dataset. This function uses the Interquartile Range (IQR) approach to detect extreme values, ensuring that any anomalies that could affect data analysis are identified. The method outputs details about outliers for each column, helping to pinpoint potential issues that may need to be addressed separately. If outliers present in a column, then it displays, if not then it displays no outliers, if the column is categorical(not numeric) the, it skips the columns

Removing Outliers

Code

```
remove_outlier <- function(dataframe, columns) {
  for (col in columns) {
    if (is.numeric(dataframe[[col]])) {
      Quantile1 <- quantile(dataframe[[col]], probs = 0.25)
      Quantile3 <- quantile(dataframe[[col]], probs = 0.75)
      IQR <- Quantile3 - Quantile1
      dataframe <- dataframe[!(
        dataframe[[col]] > Quantile3 + (IQR * 1.5) |
        dataframe[[col]] < Quantile1 - (IQR * 1.5)
      ), ]
    }
  }
}
```

```

    return(dataframe)
  }

  without_outlier_data <- remove_outlier(fresh_dataset, names(fresh_dataset))

  without_outlier_data

  detect_outlier(without_outlier_data, names(without_outlier_data))

  fresh_dataset <- without_outlier_data;

```

Output

```

> remove_outlier <- function(dataframe, columns) {
+   for (col in columns) {
+     if (is.numeric(dataframe[[col]])) {
+       quantile1 <- quantile(dataframe[[col]], probs = 0.25)
+       quantile3 <- quantile(dataframe[[col]], probs = 0.75)
+       IQR <- quantile3 - quantile1
+
+       dataframe <- dataframe[!(
+         dataframe[[col]] > quantile3 + (IQR * 1.5) |
+         dataframe[[col]] < quantile1 - (IQR * 1.5)
+       ), ]
+     }
+   }
+   return(dataframe)
+ }
> without_outlier_data <- remove_outlier(fresh_dataset, names(fresh_dataset))
> without_outlier_data
# A tibble: 176 x 14
   person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate
   <dbl> <fct> <fct> <dbl> <dbl> <fct> <dbl> <fct> <dbl>
1      21 2      3      71948      0 1      35000 1      16.0
2      25 2      1      12438      3 3       3500 3      12.9
3      23 2      2      79753      0 1      35000 3      15.2
4      24 1      3      66135      1 1      35000 3      14.3
5      21 2      1      12739      0 2       1600 4      14.7
6      22 2      1     102985      0 1      35000 4      10.4
7      21 2      4      13113      0 2       4500 5       8.83
8      23 1      2     114860      3 1      35000 4       7.9
9      23 1      2     136628      0 1      35000 6      18.2
10     24 2      3      14283      1 3       1750 2      11.0
# 166 more rows
# 5 more variables: loan_percent_income <dbl>, cb_person_cred_hist_length <dbl>, credit_score <dbl>, previous_loan_defaults_on_file <fct>,
#   loan_status <dbl>
# I use 'print(n = ...)' to see more rows
> detect_outlier(without_outlier_data, names(without_outlier_data))
No outliers detected in column person_age
column person_gender is not numeric, skipped
column person_education is not numeric, skipped
No outliers detected in column person_income
No outliers detected in column person_emp_exp
column person_home_ownership is not numeric, skipped
No outliers detected in column loan_amnt
column loan_intent is not numeric, skipped
No outliers detected in column loan_int_rate
No outliers detected in column loan_percent_income
No outliers detected in column cb_person_cred_hist_length
No outliers detected in column credit_score
column previous_loan_defaults_on_file is not numeric, skipped
No outliers detected in column loan_status
>

```

Description

We used the **remove_outlier** function to eliminate outlier values from all numeric columns in the dataset. This method applies the Interquartile Range (IQR) approach to filter out extreme values, ensuring that our dataset is free from anomalies that could skew analysis results. After removing the outliers, we re-applied the **detect_outlier** function to confirm that the dataset no longer contains any extreme values. We can see from the output that, after removing outliers, the new dataset contains 176 rows and 14 columns

Normalizing the Dataset

Code

```
normalize_dataset <- fresh_dataset;

min_age <- min(normalize_dataset$person_age, na.rm = TRUE)
max_age <- max(normalize_dataset$person_age, na.rm = TRUE)

normalize_dataset$person_age <- (normalize_dataset$person_age - min_age) / (max_age - min_age)

min_income <- min(normalize_dataset$person_income, na.rm = TRUE)
max_income <- max(normalize_dataset$person_income, na.rm = TRUE)

normalize_dataset$person_income <- (normalize_dataset$person_income - min_income) / (max_income - min_income)

min_loan_amnt <- min(normalize_dataset$loan_amnt, na.rm = TRUE)
max_loan_amnt <- max(normalize_dataset$loan_amnt, na.rm = TRUE)

normalize_dataset$loan_amnt <- (normalize_dataset$loan_amnt - min_loan_amnt) / (max_loan_amnt - min_loan_amnt);

min_loan_int_rate <- min(normalize_dataset$loan_int_rate, na.rm = TRUE)
max_loan_int_rate <- max(normalize_dataset$loan_int_rate, na.rm = TRUE)

normalize_dataset$loan_int_rate <- (normalize_dataset$loan_int_rate - min_loan_int_rate) / (max_loan_int_rate - min_loan_int_rate);

min_credit_score <- min(normalize_dataset$credit_score, na.rm = TRUE)
max_credit_score <- max(normalize_dataset$credit_score, na.rm = TRUE)

normalize_dataset$credit_score <- (normalize_dataset$credit_score - min_credit_score) / (max_credit_score - min_credit_score );

normalize_dataset
fresh_dataset <- normalize_dataset;
```

Output

```
> fresh_dataset <- without_outlier_data;
> normalize_dataset <- fresh_dataset;
> min_age <- min(normalize_dataset$person_age, na.rm = TRUE)
> max_age <- max(normalize_dataset$person_age, na.rm = TRUE)
> normalize_dataset$person_age <- (normalize_dataset$person_age - min_age) / (max_age - min_age)
> min_income <- min(normalize_dataset$person_income, na.rm = TRUE)
> max_income <- max(normalize_dataset$person_income, na.rm = TRUE)
> normalize_dataset$person_income <- (normalize_dataset$person_income - min_income) / (max_income - min_income)
> min_loan_amnt <- min(normalize_dataset$loan_amnt, na.rm = TRUE)
> max_loan_amnt <- max(normalize_dataset$loan_amnt, na.rm = TRUE)
> normalize_dataset$loan_amnt <- (normalize_dataset$loan_amnt - min_loan_amnt) / (max_loan_amnt - min_loan_amnt);
> min_loan_int_rate <- min(normalize_dataset$loan_int_rate, na.rm = TRUE)
> max_loan_int_rate <- max(normalize_dataset$loan_int_rate, na.rm = TRUE)
> normalize_dataset$loan_int_rate <- (normalize_dataset$loan_int_rate - min_loan_int_rate) / (max_loan_int_rate - min_loan_int_rate);
> min_credit_score <- min(normalize_dataset$credit_score, na.rm = TRUE)
> max_credit_score <- max(normalize_dataset$credit_score, na.rm = TRUE)
> normalize_dataset$credit_score <- (normalize_dataset$credit_score - min_credit_score) / (max_credit_score - min_credit_score);
> normalize_dataset
# A tibble: 176 x 14
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate
    <dbl> <fct> <fct> <dbl> <dbl> <fct> <dbl> <dbl> <dbl>
1 0 2 3 0.170 0 1 1 1 0.747
2 0.8 2 1 0 3 3 0.122 3 0.513
3 0.4 2 2 0.193 0 1 1 3 0.689
4 0.6 1 3 0.154 1 1 1 3 0.617
5 0 2 1 0.000862 0 2 0.00595 4 0.652
6 0.2 2 1 0.259 0 1 1 4 0.326
7 0 2 4 0.00193 0 2 0.0923 5 0.197
8 0.4 1 2 0.293 3 1 1 4 0.142
9 0.4 1 2 0.356 0 1 1 6 0.914
10 0.6 2 3 0.00528 1 3 0.0104 2 0.373
# i 166 more rows
# i 5 more variables: loan_percent_income <dbl>, cb_person_cred_hist_length <dbl>, credit_score <dbl>, previous_loan_defaults_on_file <fct>,
# loan_status <dbl>
# i Use `print(n = ...)` to see more rows
> |
```

Description

We applied **min-max normalization** to scale selected numeric columns—`person_age`, `person_income`, `loan_amnt`, and `loan_int_rate`—to a range between 0 and 1. These specific columns were chosen because they had significant variations in their data points. For instance, `person_income` had extremely high values, while `loan_amnt` and `loan_int_rate` also showed large ranges. Such disparities could skew analytical processes. Normalizing these columns helps in reducing this imbalance, ensuring all features contribute equally during model training and analysis. Other numerical columns were not prioritized as their value ranges were relatively consistent and did not pose the same scaling issues.

Descriptive Statistics

Displaying summary of the dataset

Code

```
summary(fresh_dataset);
```

Output

```
> fresh_dataset <- normalize_dataset;
> summary(fresh_dataset);
  person_age  person_gender person_education person_income  person_emp_exp  person_home_ownership  loan_amnt  loan_intent
Min.   :0.0000   1:108         1:49           Min.   :0.0000   Min.   :0.000   1:169         Min.   :0.0000   1:27
1st Qu.:0.2000   2: 68         2:63           1st Qu.:0.1382   1st Qu.:0.000   2: 3          1st Qu.:0.3147   2:49
Median :0.4000           3:21         3:21           Median :0.2085   Median :1.000   3: 3          Median :0.7024   3:22
Mean   :0.4966           4:42         Mean   :0.3532   Mean   :1.523   4: 1          Mean   :0.5648   4:29
3rd Qu.:0.8000           5: 1         3rd Qu.:0.6549   3rd Qu.:3.000   3rd Qu.:0.7775   5:21
Max.   :1.0000           Max.   :1.0000   Max.   :7.000   Max.   :1.0000   6:28
loan_int_rate  loan_percent_income  cb_person_cred_hist_length  credit_score  previous_loan_defaults_on_file  loan_status
Min.   :0.0000   Min.   :0.0000   Min.   :2.000   Min.   :0.0000   1: 46         Min.   :0.0000
1st Qu.:0.3726   1st Qu.:0.0900   1st Qu.:2.000   1st Qu.:0.4366   2:130        1st Qu.:0.0000
Median :0.4423   Median :0.2300   Median :3.000   Median :0.5986           Median :1.0000
Mean   :0.4718   Mean   :0.2293   Mean   :3.006   Mean   :0.5834           Mean   :0.6193
3rd Qu.:0.6304   3rd Qu.:0.3425   3rd Qu.:4.000   3rd Qu.:0.7477           3rd Qu.:1.0000
Max.   :1.0000   Max.   :0.5300   Max.   :4.000   Max.   :1.0000           Max.   :1.0000
> |
```

Description

By using the `summary()` function, we obtained a statistical overview of each column in the dataset, including measures like minimum, 1st quartile, median, mean, 3rd quartile, and maximum values. This helps in understanding the distribution, spread, and potential outliers across all features. It provides insights into each attribute's central tendency and variability, ensuring data integrity and highlighting areas that may need further cleaning, transformation, or normalization.

Measure of Central Tendancy

Code

```
calculate_stats <- function(dataset, columns) {
  for (column_name in columns) {
    column_data <- dataset[[column_name]]
    if (is.numeric(column_data)) {
      column_mean <- mean(column_data, na.rm = TRUE)
      column_median <- median(column_data, na.rm = TRUE)
      cat("Mean of column", column_name, "is", column_mean, "\n")
      cat("Median of column", column_name, "is", column_median, "\n")
      cat("\n")
    } else {
      column_mode <- names(sort(table(column_data), decreasing = TRUE))[1]
      cat("Mode of column", column_name, "is", column_mode, "\n")
      cat("\n")
    }
  }
}

calculate_stats(fresh_dataset, names(fresh_dataset))
```

Output

```
> calculate_stats(fresh_dataset, names(fresh_dataset))
Mean of column person_age is 0.4965909
Median of column person_age is 0.4

Mode of column person_gender is 1

Mode of column person_education is 2

Mean of column person_income is 0.3532345
Median of column person_income is 0.2084707

Mean of column person_emp_exp is 1.522727
Median of column person_emp_exp is 1

Mode of column person_home_ownership is 1

Mean of column loan_amnt is 0.5648251
Median of column loan_amnt is 0.702381

Mode of column loan_intent is 2

Mean of column loan_int_rate is 0.4718492
Median of column loan_int_rate is 0.4422504

Mean of column loan_percent_income is 0.2292614
Median of column loan_percent_income is 0.23

Mean of column cb_person_cred_hist_length is 3.005682
Median of column cb_person_cred_hist_length is 3

Mean of column credit_score is 0.5833867
Median of column credit_score is 0.5985915

Mode of column previous_loan_defaults_on_file is 2

Mean of column loan_status is 0.6193182
Median of column loan_status is 1
```

Description

We calculated mean and median for numeric columns (e.g., `person_age`, `loan_amnt`) to understand their central tendencies, while mode was calculated for categorical columns (e.g., `loan_intent`, `person_home_ownership`) to identify the most common categories. These statistics help in better understanding data distributions and ensuring that features contribute meaningfully to analysis and modeling.

Measure of spread

Code

```
columns_to_analyze <- c(
  "person_age",
  "person_income", "person_emp_exp",
  "loan_amnt", "loan_int_rate",
  "loan_percent_income", "cb_person_cred_hist_length",
  "credit_score"
)

calculate_spread <- function(dataset, columns) {
```

```

for (col_name in columns) {
  if (is.numeric(dataset[[col_name]])) {
    column_data <- dataset[[col_name]]
    column_range <- range(column_data, na.rm = TRUE)
    column_iqr <- IQR(column_data, na.rm = TRUE)
    column_sd <- sd(column_data, na.rm = TRUE)
    column_variance <- var(column_data, na.rm = TRUE)

    cat("For column", col_name, ":\n")
    cat("  Range:", column_range[2]- column_range[1], "\n")
    cat("  IQR:", column_iqr, "\n")
    cat("  Standard Deviation:", column_sd, "\n")
    cat("  Variance:", column_variance, "\n")
    cat("\n")
  }
}

calculate_spread(fresh_dataset, columns_to_analyze)

```

Output

```
> calculate_spread(fresh_dataset, columns_to_analyze)
For column person_age :
  Range: 1
  IQR: 0.6
  Standard Deviation: 0.3221176
  Variance: 0.1037597

For column person_income :
  Range: 1
  IQR: 0.5167097
  Standard Deviation: 0.3061338
  Variance: 0.09371791

For column person_emp_exp :
  Range: 7
  IQR: 3
  Standard Deviation: 1.700267
  Variance: 2.890909

For column loan_amnt :
  Range: 1
  IQR: 0.4627976
  Standard Deviation: 0.306725
  Variance: 0.0940802

For column loan_int_rate :
  Range: 1
  IQR: 0.2578241
  Standard Deviation: 0.2209408
  Variance: 0.04881482

For column loan_percent_income :
  Range: 0.53
  IQR: 0.2525
  Standard Deviation: 0.1427646
  Variance: 0.02038174

For column cb_person_cred_hist_length :
  Range: 2
  IQR: 2
  Standard Deviation: 0.7745757
  Variance: 0.5999675

For column credit_score :
  Range: 1
  IQR: 0.3110329
  Standard Deviation: 0.2130489
  Variance: 0.04538983

> |
```

Description

We analyzed the spread for selected numeric columns (e.g., person_age, loan_amnt) by calculating key metrics like range, interquartile range (IQR), standard deviation, and variance. These metrics help in understanding data dispersion, identifying potential outliers, and ensuring better model performance. Categorical columns were not analyzed in this process, as spread metrics like range or standard deviation are more meaningful for numeric data and do not apply to categorical variables.

Handling imbalance dataset

Oversampling

Code

```
class_distribution <- table(fresh_dataset$loan_status)

print(class_distribution)

if (class_distribution[1] > class_distribution[2]) {
  majority <- filter(fresh_dataset, loan_status == 0)
  minority <- filter(fresh_dataset, loan_status == 1)
} else {
  majority <- filter(fresh_dataset, loan_status == 1)
  minority <- filter(fresh_dataset, loan_status == 0)
}

set.seed(123)

oversampled_minority <- minority %>% sample_n(nrow(majority), replace = TRUE)

oversampled_data <- bind_rows(majority, oversampled_minority)

table(oversampled_data$loan_status)

oversampled_data
```

Output

```
> class_distribution <- table(fresh_dataset$loan_status)
> print(class_distribution)

 0    1 
67 109 

> if (class_distribution[1] > class_distribution[2]) {
+   majority <- filter(fresh_dataset, loan_status == 0)
+   minority <- filter(fresh_dataset, loan_status == 1)
+ } else {
+   majority <- filter(fresh_dataset, loan_status == 1)
+   minority <- filter(fresh_dataset, loan_status == 0)
+ }
> set.seed(123)
> oversampled_minority <- minority %>% sample_n(nrow(majority), replace = TRUE)
> oversampled_data <- bind_rows(majority, oversampled_minority)
> table(oversampled_data$loan_status)

 0    1 
109 109 

> oversampled_data
# A tibble: 218 x 14
   person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate
   <dbl> <fct> <fct> <dbl> <dbl> <fct> <dbl> <fct> <dbl>
1     0 2      3      0.170      0 1      1      1      0.747
2     0 2      1      0.82      0 3 3      0.122      3      0.513
3     0 2      2      0.193      0 1      1      3      0.689
4     0 1      3      0.154      1 1      1      3      0.617
5     0 2      1      0.000862 0 2      0.00595 4      0.652
6     0 2      1      0.259      0 1      1      4      0.326
7     0 2      4      0.00193 0 2      0.0923 5      0.197
8     0 4 1      0.293      3 1      1      4      0.142
9     0 4 1      0.356      0 1      1      6      0.914
10    0 6 2      0.00528 1 3      0.0104 2      0.373
# i 208 more rows
# i 5 more variables: loan_percent_income <dbl>, cb_person_cred_hist_length <dbl>, credit_score <dbl>, previous_loan_defaults_on_file <fct>,
#   loan_status <dbl>
# i Use `print(n = ...)` to see more rows
> |
```

Description

To address the class imbalance in the `loan_status` column, the majority and minority classes were first identified based on their counts. Depending on which class had more samples, it was assigned as the majority, while the other was designated as the minority. The **oversampling** technique was then applied to balance these classes by duplicating the minority class data to match the size of the majority class using the `sample_n()` function with replacement. The `bind_rows()` function was subsequently used to merge the datasets back into a single, balanced dataset.

Undersampling

Code

```
undersampled_majority <- majority %>% sample_n(nrow(minority), replace = FALSE)
undersampled_data <- bind_rows(undersampled_majority, minority)

table(undersampled_data$loan_status)

undersampled_data

fresh_dataset <- oversampled_data
```

Output

```
> undersampled_majority <- majority %>% sample_n(nrow(minority), replace = FALSE)
> undersampled_data <- bind_rows(undersampled_majority, minority)
> table(undersampled_data$loan_status)

 0  1 
67 67 
> undersampled_data
# A tibble: 134 x 14
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate
    <dbl>    <fct>        <fct>          <dbl>        <dbl>    <fct>          <dbl>    <fct>          <dbl>
1     0.4  1          2          0.293          3 1          1          4          0.142
2     0    2          2          0.00847        0 1          0.0551 6          0.722
3     0.4  1          2          0.155          0 1          0.702 2          0.400
4     0.8  2          4          0.113          3 1          0.702 2          0.772
5     0.4  1          1          0.146          0 1          0.702 1          0.539
6     0.8  1          2          0.209          3 1          0.702 2          0.687
7     0.8  1          2          0.139          0 1          0.702 5          0
8     0.4  1          1          0.139          0 1          0.702 4          0.435
9     0.2  1          3          0.339          0 1          0.792 2          0.374
10    0.8  2          4          0.731          0 1          0.702 3          0.724
# i 124 more rows
# i 5 more variables: loan_percent_income <dbl>, cb_person_cred_hist_length <dbl>, credit_score <dbl>, previous_loan_defaults_on_file <fct>,
#   loan_status <dbl>
# i Use `print(n = ...)` to see more rows
> |
```

Description

To further address class imbalance in the `loan_status` column, **undersampling** was applied. In this approach, the majority class was reduced by randomly selecting a sample equal to the size of the minority class, using the `sample_n()` function without replacement. The minority and the newly undersampled majority datasets were then combined using `bind_rows()`. This resulted in a balanced dataset where both classes were equally represented, ensuring a more stable foundation for training and analysis. The balanced dataset was stored back into `fresh_dataset`.